

# code-team-neon

August 3, 2023

## 0.1 Importing libraries

```
[54]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import skew
from scipy.stats import kurtosis
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn import linear_model
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold
from sklearn.cluster import KMeans
```

## 0.2 Reading data

```
[55]: df=pd.read_csv('heart (1).csv')
df.head()
```

```
[55]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

```
ca thal target
```

0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

```
[56]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
[57]: df.describe()
```

```
[57]:
```

	age	sex	cp	trestbps	chol \
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000
std	9.072290	0.460373	1.029641	17.516718	51.59251
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000

	fbs	restecg	thalach	exang	oldpeak \
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	0.149268	0.529756	149.114146	0.336585	1.071512
std	0.356527	0.527878	23.005724	0.472772	1.175053
min	0.000000	0.000000	71.000000	0.000000	0.000000

25%	0.000000	0.000000	132.000000	0.000000	0.000000
50%	0.000000	1.000000	152.000000	0.000000	0.800000
75%	0.000000	1.000000	166.000000	1.000000	1.800000
max	1.000000	2.000000	202.000000	1.000000	6.200000

	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

```
[58]: df.shape
```

```
[58]: (1025, 14)
```

```
[59]: df.isnull().sum()
```

```
[59]: age          0
sex            0
cp             0
trestbps      0
chol          0
fbs           0
restecg       0
thalach       0
exang         0
oldpeak       0
slope         0
ca            0
thal          0
target        0
dtype: int64
```

```
[60]: df.dtypes
```

```
[60]: age          int64
sex          int64
cp           int64
trestbps     int64
chol         int64
fbs          int64
restecg      int64
thalach      int64
```

```
exang          int64
oldpeak        float64
slope          int64
ca             int64
thal           int64
target         int64
dtype: object
```

```
[61]: df.shape
```

```
[61]: (1025, 14)
```

### 0.3 Check Outliers

```
[62]: import pandas as pd
import matplotlib.pyplot as plt

# Assuming your dataset is loaded into the DataFrame called 'df'
# Replace 'df' with your actual DataFrame name if different.

# List of columns with numeric data
numeric_columns = df.select_dtypes(include=['int64', 'float64']).columns

# Function to detect and visualize outliers using box plots
def detect_and_visualize_outliers(data_frame, column_name):
    # Create a box plot for the given column
    plt.figure(figsize=(8, 6))
    plt.boxplot(data_frame[column_name], vert=False)
    plt.title(f'Box Plot of {column_name}')
    plt.show()


    # Calculate the Interquartile Range (IQR) for the column
    Q1 = data_frame[column_name].quantile(0.25)
    Q3 = data_frame[column_name].quantile(0.75)
    IQR = Q3 - Q1

    # Calculate the lower and upper bounds for outlier detection
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

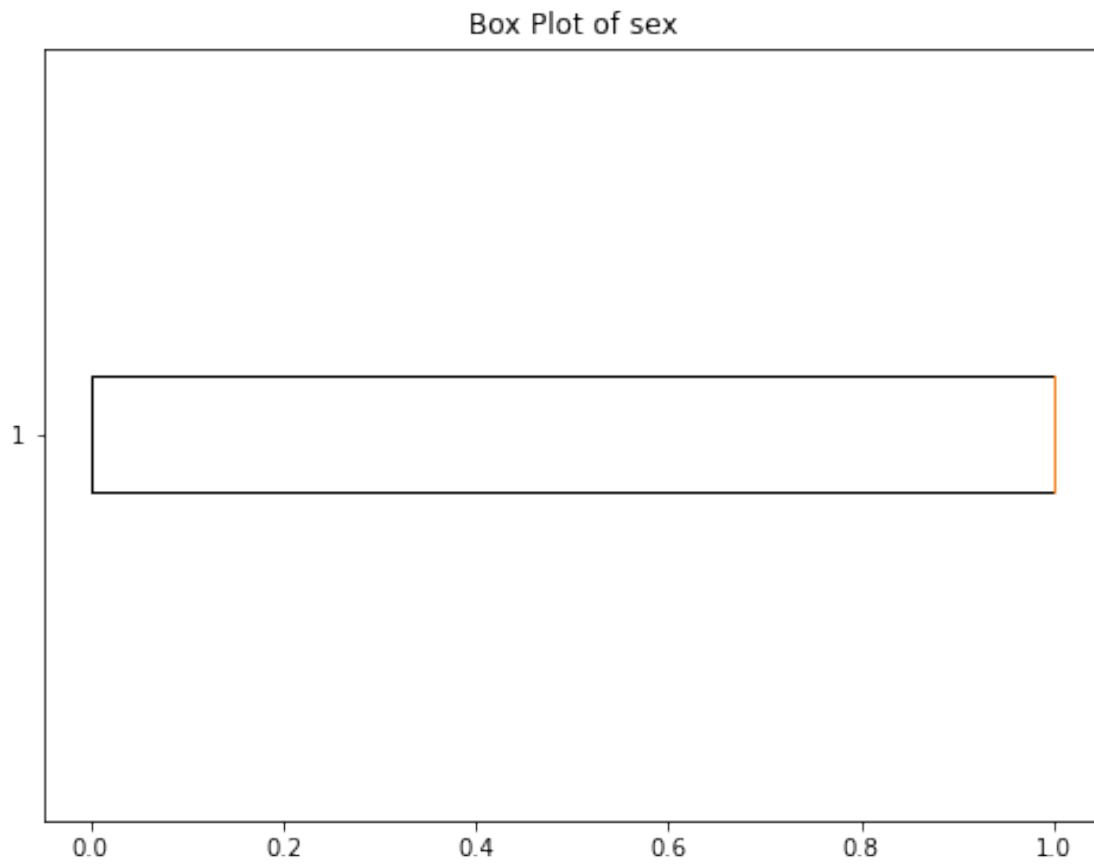
    # Find and display the outliers
    outliers = data_frame[(data_frame[column_name] < lower_bound) |
↪(data_frame[column_name] > upper_bound)]
    print(f"Outliers in '{column_name}':")
    # print(outliers)

# Loop through each numeric column and visualize outliers
```

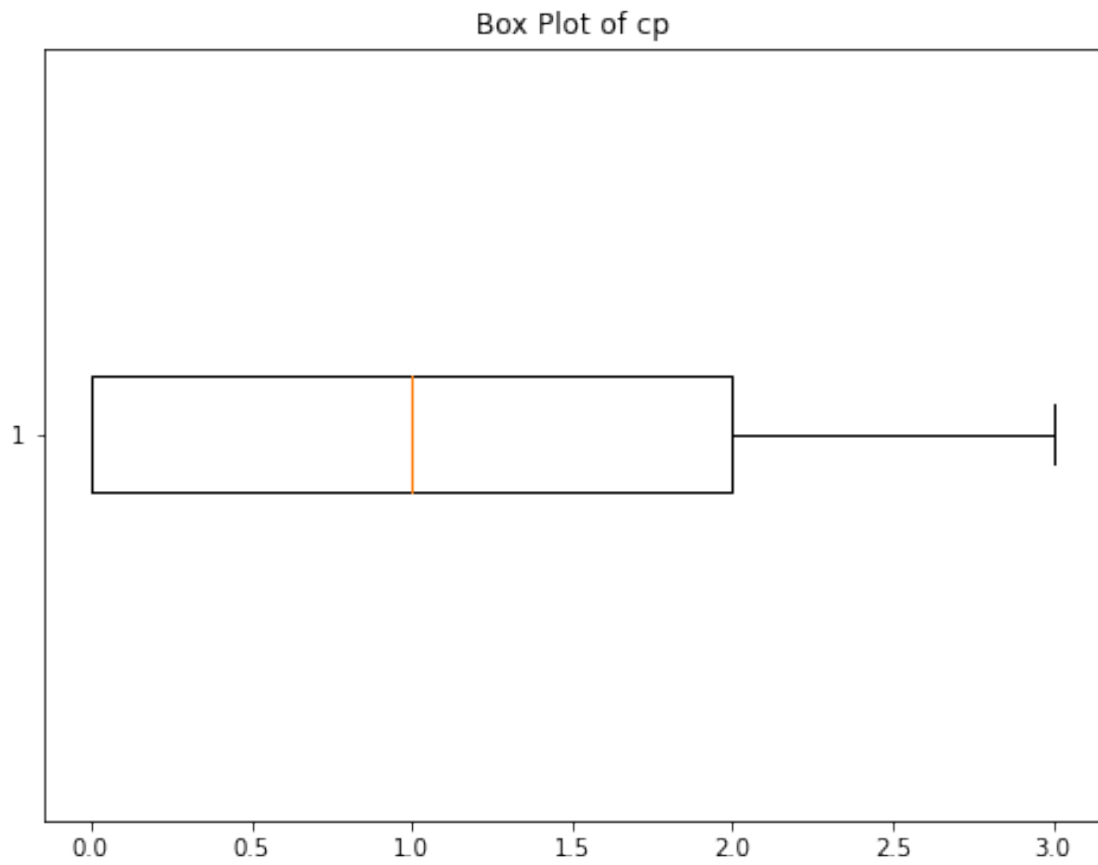
```
for col in numeric_columns:  
    detect_and_visualize_outliers(df, col)
```




Outliers in 'age':



Outliers in 'sex':

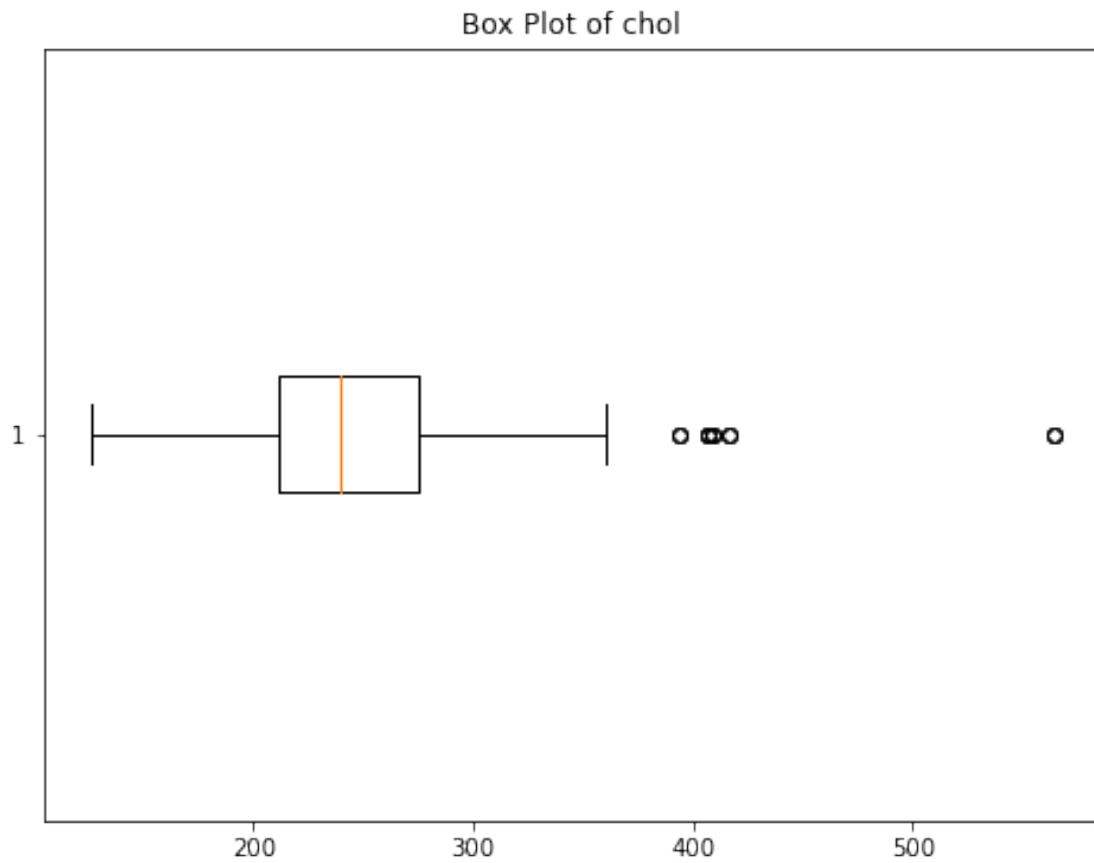


Outliers in 'cp':




Outliers in 'trestbps':

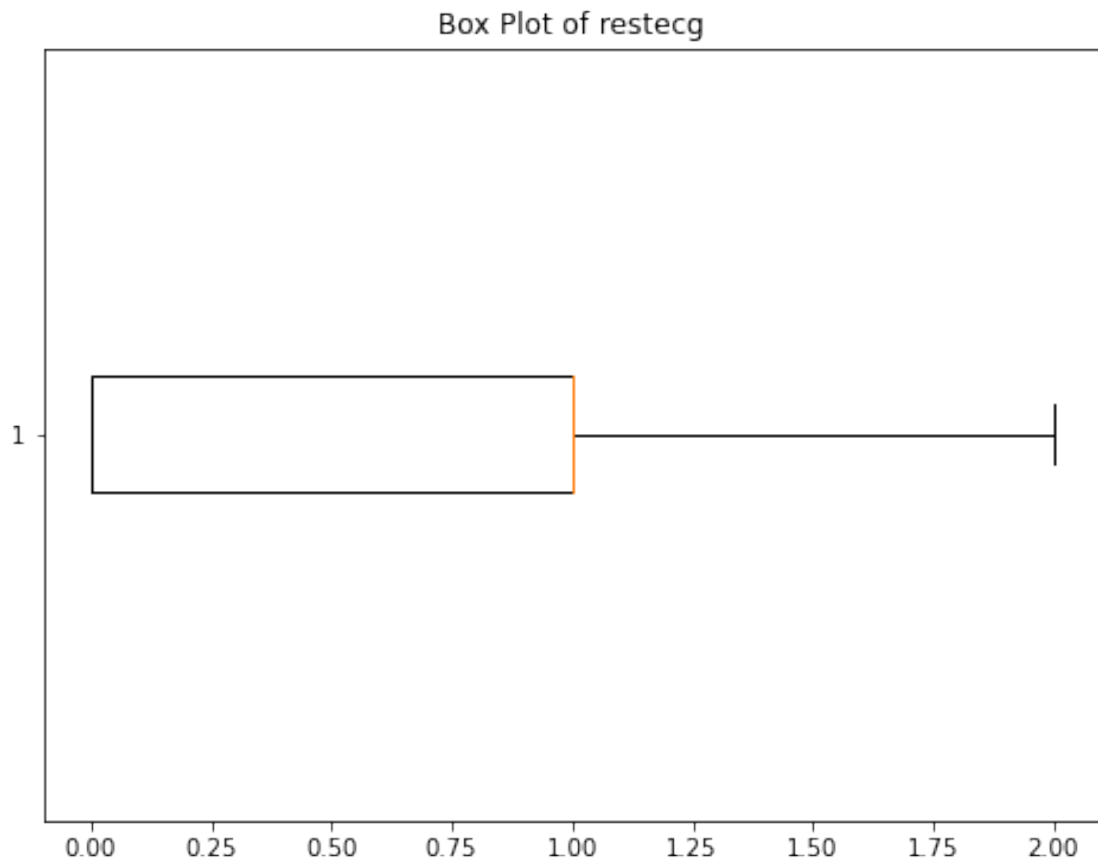




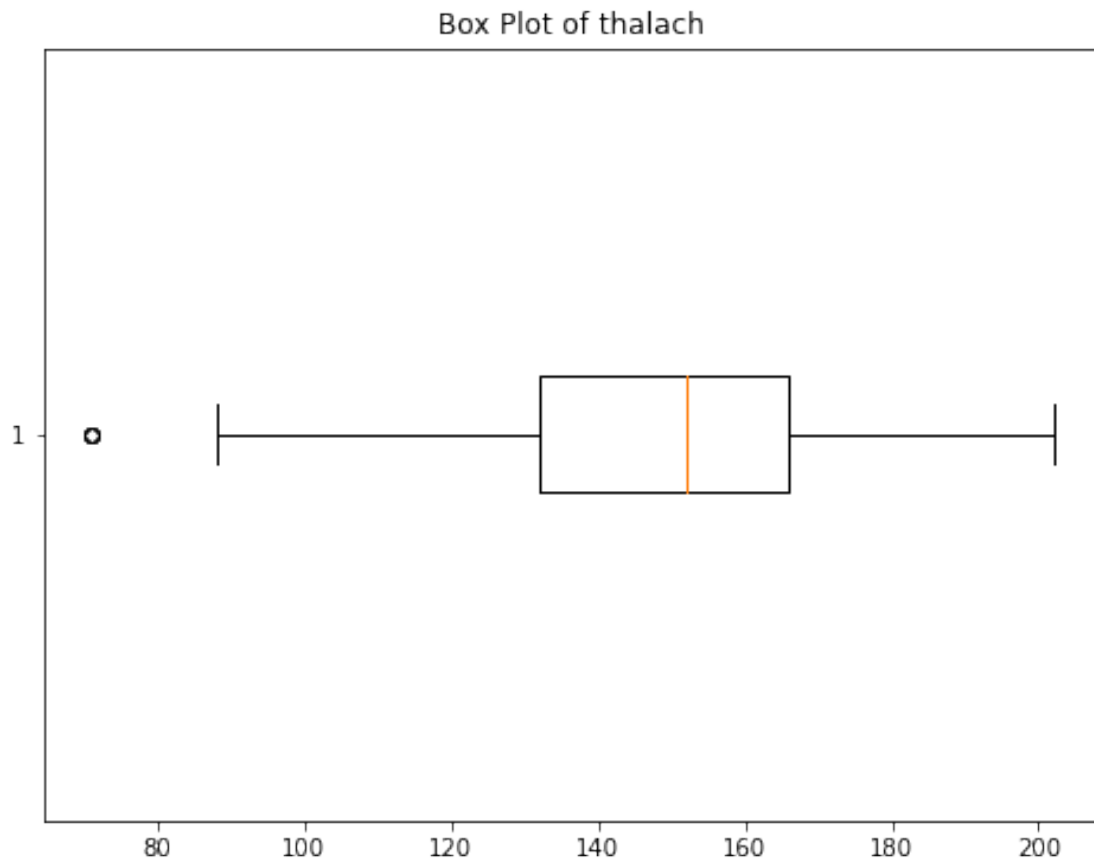
Outliers in 'chol':



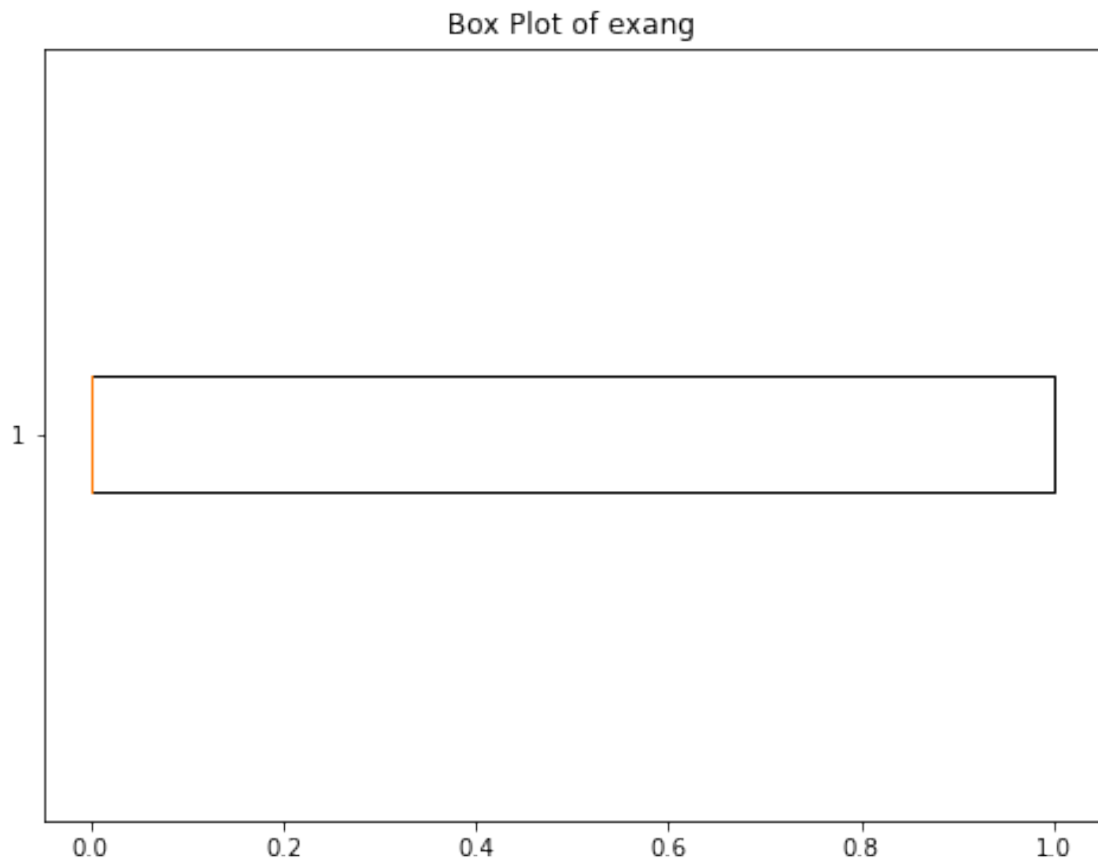
Outliers in 'fbs':




Outliers in 'restecg':



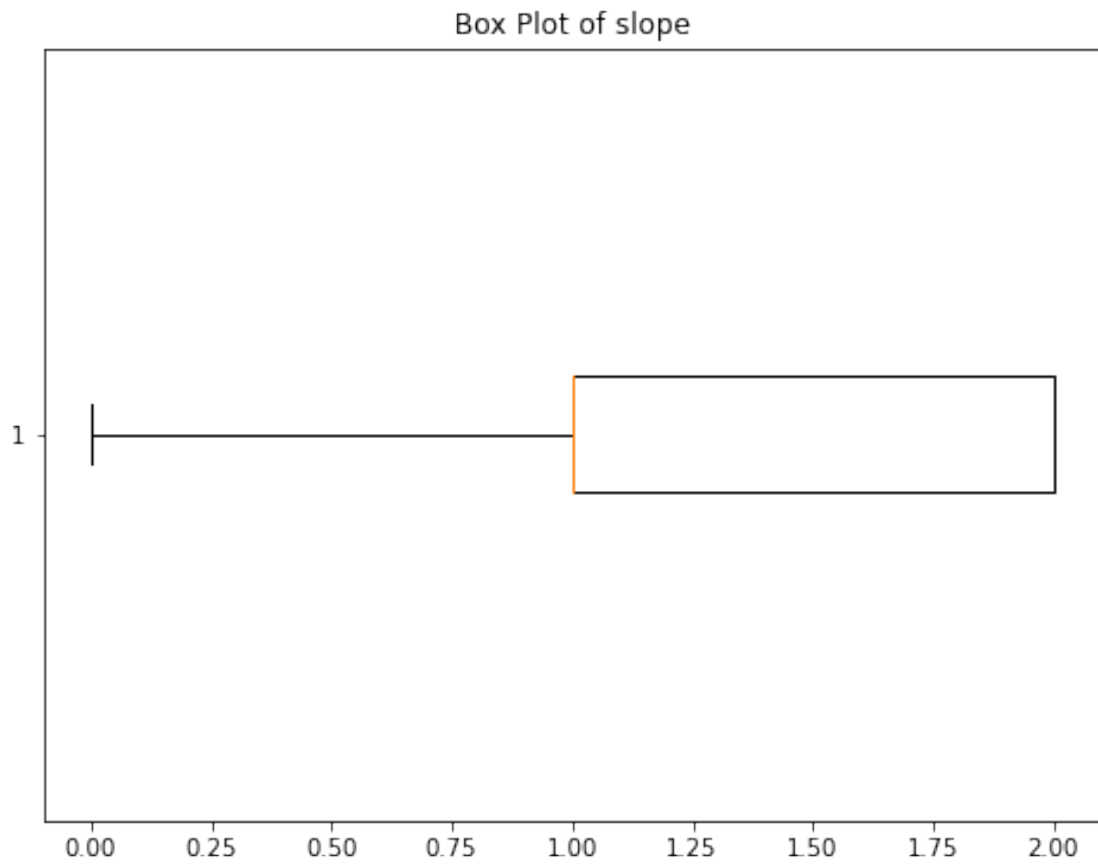
Outliers in 'thalach':




Outliers in 'exang':



Outliers in 'oldpeak':




Outliers in 'slope':

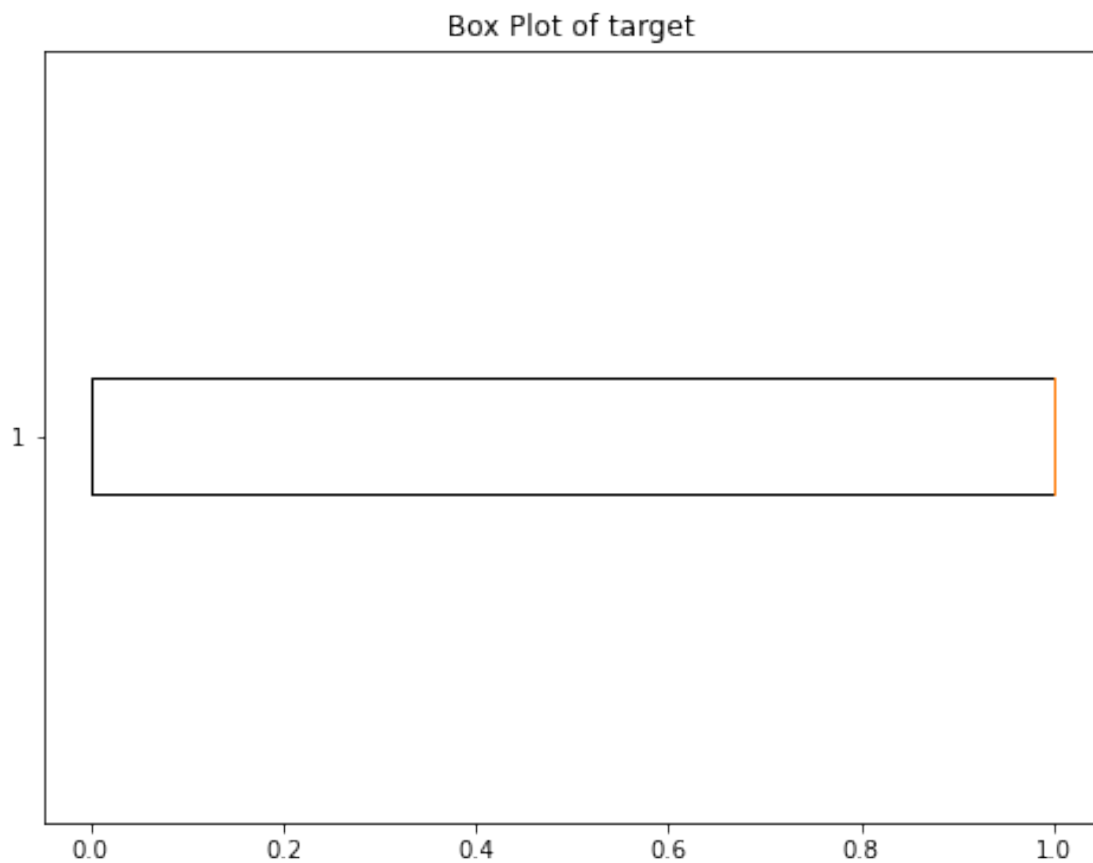


Outliers in 'ca':





Outliers in 'thal':



Outliers in 'target':

## 0.4 Data Split

```
[63]: x=df.drop('target',axis=1)
      x.head()
```

```
[63]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

	ca	thal
0	2	3
1	0	3
2	0	3
3	1	3
4	3	2

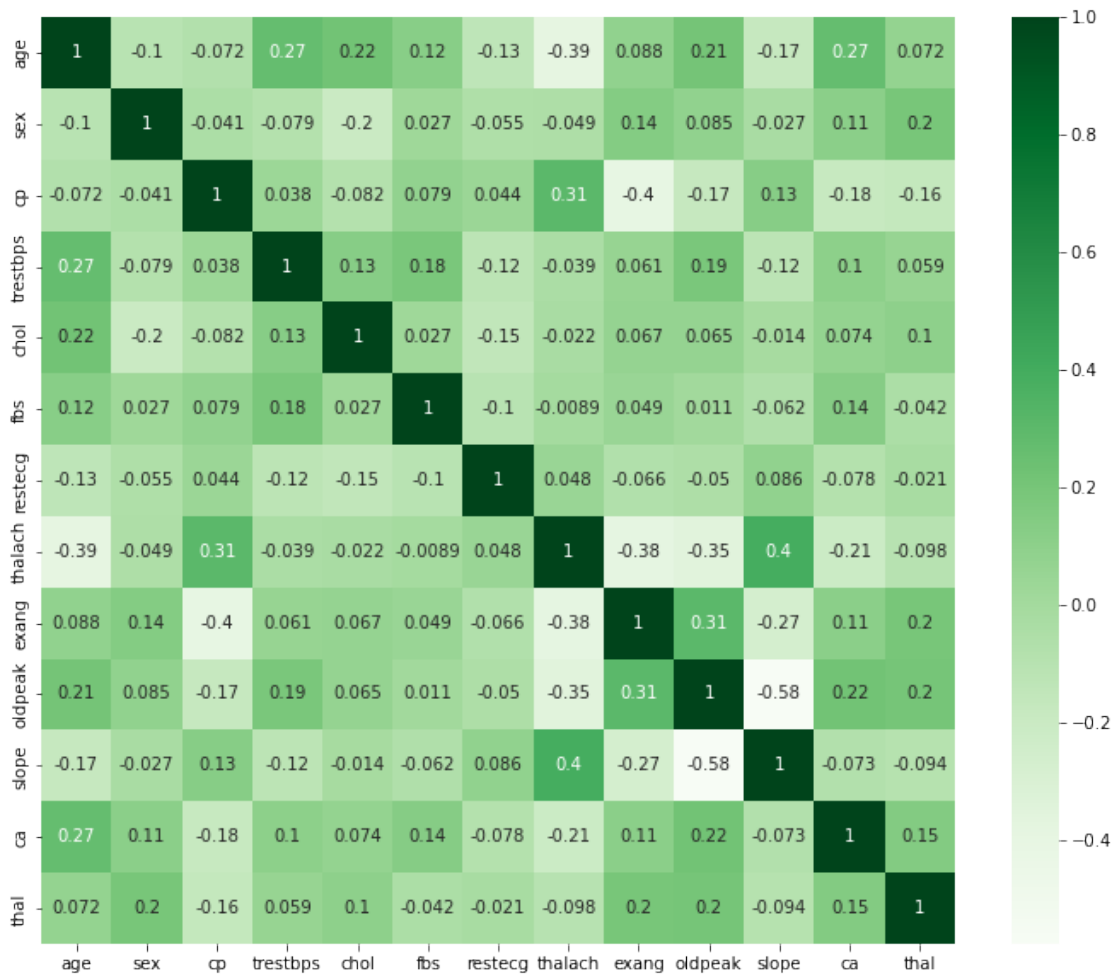
```
[64]: y=df.target
      y.head()
```

```
[64]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: target, dtype: int64
```

```
[65]: X_train, X_test, y_train, y_test =train_test_split(x,y,test_size=0.25)
```

## 1 Correlation Matrix

```
[66]: plt.figure(figsize=(12,10))
      cor=x.corr()
      sns.heatmap(cor,annot=True,cmap=plt.cm.Greens)
      plt.show()
```



## 2 Neural Network

```
[67]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import Adam
from keras.callbacks import EarlyStopping

# Step 1: Load the dataset
# Replace this with your actual dataset loading process
# Assuming 'df' contains the dataset with the features and target 'cardio'
# For example:
# df = pd.read_csv('your_dataset.csv')

# Step 2: Data Preprocessing
# Perform data preprocessing steps here, such as encoding categorical variables,
# and scaling
X = df.drop(columns=['target'])
y = df['target']

# Perform feature scaling using StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.
# 25, random_state=42)

# Step 4: Build the Neural Network Model
model = Sequential()
model.add(Dense(64, input_dim=X_train.shape[1], activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

# Step 5: Compile the model
optimizer = Adam(learning_rate=0.001)
model.compile(loss='binary_crossentropy', optimizer=optimizer,
# metrics=['accuracy'])

# Step 6: Train the model
early_stopping = EarlyStopping(patience=5, restore_best_weights=True)
```

```
history = model.fit(X_train, y_train, validation_split=0.2, epochs=100,  
    ↪batch_size=64, callbacks=[early_stopping])
```

```
# Step 7: Evaluate the model on the test set
```

```
test_loss, test_accuracy = model.evaluate(X_test, y_test)
```

```
print("Test Loss:", test_loss)
```

```
print("Test Accuracy:", test_accuracy)
```

Epoch 1/100

10/10 [=====] - 1s 22ms/step - loss: 0.7826 - accuracy:  
0.3420 - val\_loss: 0.6932 - val\_accuracy: 0.5455

Epoch 2/100

10/10 [=====] - 0s 5ms/step - loss: 0.6276 - accuracy:  
0.6954 - val\_loss: 0.6058 - val\_accuracy: 0.7143

Epoch 3/100

10/10 [=====] - 0s 4ms/step - loss: 0.5222 - accuracy:  
0.8241 - val\_loss: 0.5423 - val\_accuracy: 0.7857

Epoch 4/100

10/10 [=====] - 0s 4ms/step - loss: 0.4462 - accuracy:  
0.8550 - val\_loss: 0.4964 - val\_accuracy: 0.8052

Epoch 5/100

10/10 [=====] - 0s 5ms/step - loss: 0.3914 - accuracy:  
0.8664 - val\_loss: 0.4696 - val\_accuracy: 0.8052

Epoch 6/100

10/10 [=====] - 0s 5ms/step - loss: 0.3531 - accuracy:  
0.8632 - val\_loss: 0.4528 - val\_accuracy: 0.8052

Epoch 7/100

10/10 [=====] - 0s 5ms/step - loss: 0.3292 - accuracy:  
0.8697 - val\_loss: 0.4408 - val\_accuracy: 0.8247

Epoch 8/100

10/10 [=====] - 0s 5ms/step - loss: 0.3113 - accuracy:  
0.8827 - val\_loss: 0.4375 - val\_accuracy: 0.8117

Epoch 9/100

10/10 [=====] - 0s 5ms/step - loss: 0.2973 - accuracy:  
0.8876 - val\_loss: 0.4351 - val\_accuracy: 0.8182

Epoch 10/100

10/10 [=====] - 0s 5ms/step - loss: 0.2857 - accuracy:  
0.8941 - val\_loss: 0.4278 - val\_accuracy: 0.8117

Epoch 11/100

10/10 [=====] - 0s 5ms/step - loss: 0.2760 - accuracy:  
0.8958 - val\_loss: 0.4224 - val\_accuracy: 0.8117

Epoch 12/100

10/10 [=====] - 0s 5ms/step - loss: 0.2670 - accuracy:  
0.8990 - val\_loss: 0.4180 - val\_accuracy: 0.8117

Epoch 13/100

10/10 [=====] - 0s 5ms/step - loss: 0.2591 - accuracy:  
0.9007 - val\_loss: 0.4166 - val\_accuracy: 0.8247

Epoch 14/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2528 - accuracy: 0.9055 - val\_loss: 0.4116 - val\_accuracy: 0.8247  
Epoch 15/100  
10/10 [=====] - 0s 4ms/step - loss: 0.2458 - accuracy: 0.9055 - val\_loss: 0.4101 - val\_accuracy: 0.8312  
Epoch 16/100  
10/10 [=====] - 0s 6ms/step - loss: 0.2395 - accuracy: 0.9072 - val\_loss: 0.4038 - val\_accuracy: 0.8312  
Epoch 17/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2348 - accuracy: 0.9153 - val\_loss: 0.4007 - val\_accuracy: 0.8377  
Epoch 18/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2285 - accuracy: 0.9169 - val\_loss: 0.4007 - val\_accuracy: 0.8377  
Epoch 19/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2226 - accuracy: 0.9202 - val\_loss: 0.3966 - val\_accuracy: 0.8377  
Epoch 20/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2179 - accuracy: 0.9218 - val\_loss: 0.3944 - val\_accuracy: 0.8312  
Epoch 21/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2113 - accuracy: 0.9202 - val\_loss: 0.3912 - val\_accuracy: 0.8377  
Epoch 22/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2059 - accuracy: 0.9218 - val\_loss: 0.3902 - val\_accuracy: 0.8377  
Epoch 23/100  
10/10 [=====] - 0s 5ms/step - loss: 0.2007 - accuracy: 0.9202 - val\_loss: 0.3869 - val\_accuracy: 0.8377  
Epoch 24/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1961 - accuracy: 0.9202 - val\_loss: 0.3827 - val\_accuracy: 0.8442  
Epoch 25/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1910 - accuracy: 0.9251 - val\_loss: 0.3808 - val\_accuracy: 0.8442  
Epoch 26/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1857 - accuracy: 0.9235 - val\_loss: 0.3779 - val\_accuracy: 0.8442  
Epoch 27/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1809 - accuracy: 0.9251 - val\_loss: 0.3762 - val\_accuracy: 0.8442  
Epoch 28/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1765 - accuracy: 0.9251 - val\_loss: 0.3761 - val\_accuracy: 0.8442  
Epoch 29/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1726 - accuracy: 0.9365 - val\_loss: 0.3743 - val\_accuracy: 0.8442

Epoch 30/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1674 - accuracy: 0.9365 - val\_loss: 0.3720 - val\_accuracy: 0.8442

Epoch 31/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1629 - accuracy: 0.9365 - val\_loss: 0.3697 - val\_accuracy: 0.8571

Epoch 32/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1583 - accuracy: 0.9446 - val\_loss: 0.3684 - val\_accuracy: 0.8571

Epoch 33/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1539 - accuracy: 0.9463 - val\_loss: 0.3652 - val\_accuracy: 0.8571

Epoch 34/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1491 - accuracy: 0.9430 - val\_loss: 0.3629 - val\_accuracy: 0.8571

Epoch 35/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1448 - accuracy: 0.9463 - val\_loss: 0.3537 - val\_accuracy: 0.8571

Epoch 36/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1401 - accuracy: 0.9479 - val\_loss: 0.3525 - val\_accuracy: 0.8571

Epoch 37/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1353 - accuracy: 0.9528 - val\_loss: 0.3519 - val\_accuracy: 0.8571

Epoch 38/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1324 - accuracy: 0.9593 - val\_loss: 0.3472 - val\_accuracy: 0.8701

Epoch 39/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1274 - accuracy: 0.9658 - val\_loss: 0.3473 - val\_accuracy: 0.8701

Epoch 40/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1246 - accuracy: 0.9577 - val\_loss: 0.3480 - val\_accuracy: 0.8636

Epoch 41/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1203 - accuracy: 0.9642 - val\_loss: 0.3377 - val\_accuracy: 0.8701

Epoch 42/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1158 - accuracy: 0.9707 - val\_loss: 0.3365 - val\_accuracy: 0.8896

Epoch 43/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1129 - accuracy: 0.9739 - val\_loss: 0.3374 - val\_accuracy: 0.8896

Epoch 44/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1084 - accuracy: 0.9739 - val\_loss: 0.3297 - val\_accuracy: 0.8896

Epoch 45/100  
10/10 [=====] - 0s 4ms/step - loss: 0.1050 - accuracy: 0.9739 - val\_loss: 0.3335 - val\_accuracy: 0.8896

Epoch 46/100  
10/10 [=====] - 0s 5ms/step - loss: 0.1010 - accuracy:  
0.9739 - val\_loss: 0.3312 - val\_accuracy: 0.8896  
Epoch 47/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0974 - accuracy:  
0.9788 - val\_loss: 0.3270 - val\_accuracy: 0.8896  
Epoch 48/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0942 - accuracy:  
0.9805 - val\_loss: 0.3249 - val\_accuracy: 0.8896  
Epoch 49/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0915 - accuracy:  
0.9821 - val\_loss: 0.3199 - val\_accuracy: 0.8961  
Epoch 50/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0884 - accuracy:  
0.9805 - val\_loss: 0.3170 - val\_accuracy: 0.8961  
Epoch 51/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0865 - accuracy:  
0.9870 - val\_loss: 0.3258 - val\_accuracy: 0.9026  
Epoch 52/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0822 - accuracy:  
0.9870 - val\_loss: 0.3154 - val\_accuracy: 0.8961  
Epoch 53/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0803 - accuracy:  
0.9870 - val\_loss: 0.3135 - val\_accuracy: 0.9026  
Epoch 54/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0766 - accuracy:  
0.9886 - val\_loss: 0.3152 - val\_accuracy: 0.9091  
Epoch 55/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0746 - accuracy:  
0.9886 - val\_loss: 0.3098 - val\_accuracy: 0.9091  
Epoch 56/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0715 - accuracy:  
0.9886 - val\_loss: 0.3092 - val\_accuracy: 0.9091  
Epoch 57/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0689 - accuracy:  
0.9886 - val\_loss: 0.3089 - val\_accuracy: 0.9091  
Epoch 58/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0670 - accuracy:  
0.9886 - val\_loss: 0.3071 - val\_accuracy: 0.9091  
Epoch 59/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0647 - accuracy:  
0.9886 - val\_loss: 0.3065 - val\_accuracy: 0.9091  
Epoch 60/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0625 - accuracy:  
0.9886 - val\_loss: 0.2994 - val\_accuracy: 0.9091  
Epoch 61/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0606 - accuracy:  
0.9919 - val\_loss: 0.2962 - val\_accuracy: 0.9091



Epoch 62/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0583 - accuracy:  
0.9886 - val\_loss: 0.2963 - val\_accuracy: 0.9091  
Epoch 63/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0563 - accuracy:  
0.9886 - val\_loss: 0.2953 - val\_accuracy: 0.9091  
Epoch 64/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0548 - accuracy:  
0.9951 - val\_loss: 0.2939 - val\_accuracy: 0.9091  
Epoch 65/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0534 - accuracy:  
0.9935 - val\_loss: 0.2940 - val\_accuracy: 0.9156  
Epoch 66/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0515 - accuracy:  
0.9951 - val\_loss: 0.2921 - val\_accuracy: 0.9091  
Epoch 67/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0501 - accuracy:  
0.9951 - val\_loss: 0.2909 - val\_accuracy: 0.9156  
Epoch 68/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0480 - accuracy:  
0.9967 - val\_loss: 0.2866 - val\_accuracy: 0.9156  
Epoch 69/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0468 - accuracy:  
0.9967 - val\_loss: 0.2897 - val\_accuracy: 0.9156  
Epoch 70/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0453 - accuracy:  
0.9967 - val\_loss: 0.2849 - val\_accuracy: 0.9156  
Epoch 71/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0448 - accuracy:  
0.9967 - val\_loss: 0.2881 - val\_accuracy: 0.9156  
Epoch 72/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0430 - accuracy:  
0.9967 - val\_loss: 0.2883 - val\_accuracy: 0.9156  
Epoch 73/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0416 - accuracy:  
0.9967 - val\_loss: 0.2846 - val\_accuracy: 0.9156  
Epoch 74/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0401 - accuracy:  
0.9967 - val\_loss: 0.2781 - val\_accuracy: 0.9156  
Epoch 75/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0388 - accuracy:  
0.9967 - val\_loss: 0.2807 - val\_accuracy: 0.9156  
Epoch 76/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0376 - accuracy:  
0.9967 - val\_loss: 0.2771 - val\_accuracy: 0.9156  
Epoch 77/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0367 - accuracy:  
0.9967 - val\_loss: 0.2819 - val\_accuracy: 0.9156

Epoch 78/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0354 - accuracy:  
0.9967 - val\_loss: 0.2790 - val\_accuracy: 0.9156  
Epoch 79/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0345 - accuracy:  
0.9967 - val\_loss: 0.2821 - val\_accuracy: 0.9156  
Epoch 80/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0337 - accuracy:  
0.9967 - val\_loss: 0.2763 - val\_accuracy: 0.9156  
Epoch 81/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0326 - accuracy:  
0.9967 - val\_loss: 0.2788 - val\_accuracy: 0.9156  
Epoch 82/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0316 - accuracy:  
0.9967 - val\_loss: 0.2778 - val\_accuracy: 0.9156  
Epoch 83/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0308 - accuracy:  
0.9967 - val\_loss: 0.2804 - val\_accuracy: 0.9156  
Epoch 84/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0300 - accuracy:  
0.9967 - val\_loss: 0.2755 - val\_accuracy: 0.9156  
Epoch 85/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0295 - accuracy:  
0.9967 - val\_loss: 0.2776 - val\_accuracy: 0.9156  
Epoch 86/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0283 - accuracy:  
0.9967 - val\_loss: 0.2745 - val\_accuracy: 0.9156  
Epoch 87/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0276 - accuracy:  
0.9967 - val\_loss: 0.2726 - val\_accuracy: 0.9286  
Epoch 88/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0269 - accuracy:  
0.9967 - val\_loss: 0.2706 - val\_accuracy: 0.9286  
Epoch 89/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0263 - accuracy:  
0.9967 - val\_loss: 0.2728 - val\_accuracy: 0.9286  
Epoch 90/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0255 - accuracy:  
0.9967 - val\_loss: 0.2736 - val\_accuracy: 0.9286  
Epoch 91/100  
10/10 [=====] - 0s 4ms/step - loss: 0.0248 - accuracy:  
0.9967 - val\_loss: 0.2707 - val\_accuracy: 0.9286  
Epoch 92/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0241 - accuracy:  
0.9967 - val\_loss: 0.2731 - val\_accuracy: 0.9286  
Epoch 93/100  
10/10 [=====] - 0s 5ms/step - loss: 0.0236 - accuracy:  
0.9967 - val\_loss: 0.2702 - val\_accuracy: 0.9286

```

Epoch 94/100
10/10 [=====] - 0s 5ms/step - loss: 0.0227 - accuracy:
0.9967 - val_loss: 0.2712 - val_accuracy: 0.9286
Epoch 95/100
10/10 [=====] - 0s 5ms/step - loss: 0.0222 - accuracy:
0.9967 - val_loss: 0.2701 - val_accuracy: 0.9286
Epoch 96/100
10/10 [=====] - 0s 5ms/step - loss: 0.0217 - accuracy:
0.9967 - val_loss: 0.2693 - val_accuracy: 0.9286
Epoch 97/100
10/10 [=====] - 0s 5ms/step - loss: 0.0212 - accuracy:
0.9967 - val_loss: 0.2699 - val_accuracy: 0.9286
Epoch 98/100
10/10 [=====] - 0s 5ms/step - loss: 0.0207 - accuracy:
0.9967 - val_loss: 0.2719 - val_accuracy: 0.9286
Epoch 99/100
10/10 [=====] - 0s 5ms/step - loss: 0.0202 - accuracy:
0.9967 - val_loss: 0.2719 - val_accuracy: 0.9416
Epoch 100/100
10/10 [=====] - 0s 5ms/step - loss: 0.0194 - accuracy:
0.9967 - val_loss: 0.2678 - val_accuracy: 0.9286
9/9 [=====] - 0s 2ms/step - loss: 0.2137 - accuracy:
0.9572
Test Loss: 0.21372650563716888
Test Accuracy: 0.957198441028595

```

```


[68]: import matplotlib.pyplot as plt

# Access the training accuracy and validation accuracy from the history object
train_accuracy = history.history['accuracy']
val_accuracy = history.history['val_accuracy']

# Create a list with the number of epochs
epochs = range(1, len(train_accuracy) + 1)

# Plot the training accuracy and validation accuracy over epochs
plt.figure(figsize=(8, 6))
plt.plot(epochs, train_accuracy, label='Training Accuracy', marker='o')
plt.plot(epochs, val_accuracy, label='Validation Accuracy', marker='o')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.xticks(epochs)
plt.show()

```




```
[69]: import matplotlib.pyplot as plt

# Access the training loss and validation loss from the history object
train_loss = history.history['loss']
val_loss = history.history['val_loss']

# Create a list with the number of epochs
epochs = range(1, len(train_loss) + 1)

# Plot the training loss and validation loss over epochs
plt.figure(figsize=(8, 6))
plt.plot(epochs, train_loss, label='Training Loss', marker='o')
plt.plot(epochs, val_loss, label='Validation Loss', marker='o')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.xticks(epochs)
plt.show()
```



### 3 Logistic Regression model

```
[70]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, \
    ↪confusion_matrix

# Separate features (independent variables) and target variable (dependent
    ↪variable)
X = df.drop(columns=['target'])
y = df['target']

# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
    ↪random_state=42)

# Initialize the Logistic Regression model
```

```

model = LogisticRegression()

# Train the model on the training data
model.fit(X_train, y_train)

# Predict the target values on the test data
y_pred = model.predict(X_test)

# Calculate accuracy on the test data
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

# Generate the classification report
report = classification_report(y_test, y_pred)
print("Classification Report:")
print(report)

# Generate the confusion matrix
confusion = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(confusion)

```

Accuracy: 0.7804878048780488

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.69	0.76	102
1	0.74	0.87	0.80	103
accuracy			0.78	205
macro avg	0.79	0.78	0.78	205
weighted avg	0.79	0.78	0.78	205

Confusion Matrix:

```

[[70 32]
 [13 90]]

```

C:\Users\gayathriboddu\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

## 4 GradientBoostingClassifier

```
[71]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, confusion_matrix

X = df.drop(columns=['target'])
y = df['target']
param_grid = {
    'n_estimators': [50, 100, 150],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 4, 5]
}

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    ↪random_state=42)
model = GradientBoostingClassifier(random_state=42)
grid_search = GridSearchCV(model, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Get the best hyperparameters
best_model = grid_search.best_estimator_
best_model.fit(X_train, y_train)
train_accuracy = best_model.score(X_train, y_train)
test_accuracy = best_model.score(X_test, y_test)

print("Train Accuracy:", train_accuracy)
print("Test Accuracy:", test_accuracy)
y_pred = best_model.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)
```

```
Train Accuracy: 1.0
Test Accuracy: 0.9766536964980544
Confusion Matrix:
[[132  0]
 [ 6 119]]
```

## 5 SVM

```
[73]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix

# Replace 'target' with the actual column name for the target variable in your \
    dataset
target_column_name = 'target'

# Separate features (independent variables) and target variable (dependent \
    variable)
X = df.drop(columns=[target_column_name])
y = df[target_column_name]

# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
    random_state=42)

# Initialize the SVM classifier with a linear kernel
svm_classifier = SVC(kernel='linear')

# Train the classifier on the training data
svm_classifier.fit(X_train, y_train)

# Predict the target values on the test data
y_pred = svm_classifier.predict(X_test)

# Calculate accuracy on the test data
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

# Generate the classification report
report = classification_report(y_test, y_pred)
print("Classification Report:")
print(report)

# Generate the confusion matrix
confusion = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(confusion)
```

Accuracy: 0.8048780487804879

Classification Report:



	precision	recall	f1-score	support
0	0.88	0.71	0.78	102
1	0.76	0.90	0.82	103
accuracy			0.80	205
macro avg	0.82	0.80	0.80	205
weighted avg	0.82	0.80	0.80	205

Confusion Matrix:

```
[[72 30]
 [10 93]]
```

## 6 k-NN classifier

```
[74]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix

# Separate features (independent variables) and target variable (dependent
    variable)
X = df.drop(columns=['target'])
y = df['target']

# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
    random_state=42)

# Initialize the k-NN classifier with k=3 (you can choose any value for k)
knn = KNeighborsClassifier(n_neighbors=3)

# Train the classifier on the training data
knn.fit(X_train, y_train)

# Predict the target values on the test data
y_pred = knn.predict(X_test)

# Calculate accuracy on the test data
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

# Generate the classification report
report = classification_report(y_test, y_pred)
print("Classification Report:")
```

```

print(report)

# Generate the confusion matrix
confusion = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(confusion)

```

Accuracy: 0.9024390243902439

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.89	0.90	102
1	0.90	0.91	0.90	103
accuracy			0.90	205
macro avg	0.90	0.90	0.90	205
weighted avg	0.90	0.90	0.90	205

Confusion Matrix:

```

[[91 11]
 [ 9 94]]

```

## 7 Naive Bayes classifier

```

[75]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix

# Separate features (independent variables) and target variable (dependent
    variable)
X = df.drop(columns=['target'])
y = df['target']

# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
    random_state=42)

# Initialize the Naive Bayes classifier (Gaussian Naive Bayes)
naive_bayes_classifier = GaussianNB()

# Train the classifier on the training data
naive_bayes_classifier.fit(X_train, y_train)

```

```

# Predict the target values on the test data
y_pred = naive_bayes_classifier.predict(X_test)

# Calculate accuracy on the test data
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

# Generate the classification report
report = classification_report(y_test, y_pred)
print("Classification Report:")
print(report)

# Generate the confusion matrix
confusion = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(confusion)

```

Accuracy: 0.8

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.71	0.78	102
1	0.75	0.89	0.82	103
accuracy			0.80	205
macro avg	0.81	0.80	0.80	205
weighted avg	0.81	0.80	0.80	205

Confusion Matrix:

```

[[72 30]
 [11 92]]

```

## 8 Train the Gradient Boosting Classifier with Regularization

```

[39]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix

# Step 1: Load the dataset
# Replace this with your actual dataset loading process
# Assuming 'df' contains the dataset with the features and target 'cardio'
# For example:
# df = pd.read_csv('your_dataset.csv')

```

```

# Step 2: Data Preprocessing (if needed)
# If needed, perform data preprocessing steps here, such as encoding,
↳ categorical variables, scaling, etc.

# Step 3: Train-Test Split
X = df.drop(columns=['target'])
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳ random_state=42)

# Step 4: Train the Gradient Boosting Classifier with Regularization
# Apply L1 (Lasso) regularization
model_l1 = XGBClassifier(booster='gbtree', reg_alpha=0.1, random_state=42)
model_l1.fit(X_train, y_train)

# Apply L2 (Ridge) regularization
model_l2 = XGBClassifier(booster='gbtree', reg_lambda=0.1, random_state=42)
model_l2.fit(X_train, y_train)

# Step 5: Evaluate the models
y_pred_l1 = model_l1.predict(X_test)
y_pred_l2 = model_l2.predict(X_test)

accuracy_l1 = accuracy_score(y_test, y_pred_l1)
accuracy_l2 = accuracy_score(y_test, y_pred_l2)

print("Accuracy with L1 Regularization:", accuracy_l1)
print("Accuracy with L2 Regularization:", accuracy_l2)

# Step 6: Calculate and compare the confusion matrices
conf_matrix_l1 = confusion_matrix(y_test, y_pred_l1)
conf_matrix_l2 = confusion_matrix(y_test, y_pred_l2)

print("Confusion Matrix with L1 Regularization:")
print(conf_matrix_l1)

print("Confusion Matrix with L2 Regularization:")
print(conf_matrix_l2)

```

Accuracy with L1 Regularization: 0.9883268482490273

Accuracy with L2 Regularization: 0.9883268482490273

Confusion Matrix with L1 Regularization:

```
[[132  0]
 [ 3 122]]
```

Confusion Matrix with L2 Regularization:

```
[[132  0]
```

```
[ 3 122]]
```

## 9 Random Forest classifier

```
[40]: import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report

# Load the dataset and perform train-test split (replace with your data loading \
    process)
# Assuming 'df' contains the dataset with the features and target 'target'
# For example:
# df = pd.read_csv('your_dataset.csv')

X = df.drop(columns=['target'])
y = df['target']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, \
    random_state=42)

# Create a Random Forest classifier
rf_classifier = RandomForestClassifier(random_state=42)

# Perform GridSearchCV to find the best hyperparameters
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

grid_search = GridSearchCV(rf_classifier, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Get the best hyperparameters
best_rf_model = grid_search.best_estimator_

# Fit the best model to the training data
best_rf_model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = best_rf_model.predict(X_test)

# Calculate accuracy
```

```

test_accuracy = accuracy_score(y_test, y_pred)

# Print the results
print("Best Hyperparameters:", grid_search.best_params_)
print("Test Accuracy:", test_accuracy)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

# Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Best Hyperparameters: {'max\_depth': 10, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 50}  
Test Accuracy: 0.9883268482490273

Confusion Matrix:

```
[[132  0]
 [ 3 122]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	132
1	1.00	0.98	0.99	125
accuracy			0.99	257
macro avg	0.99	0.99	0.99	257
weighted avg	0.99	0.99	0.99	257

## 10 Decision Tree Classifier

```

[41]: import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,
    classification_report

# Load the dataset and perform train-test split (replace with your data loading
    process)
# Assuming 'df' contains the dataset with the features and target 'target'
# For example:
df = pd.read_csv('your_dataset.csv')

```

```

X = df.drop(columns=['target'])
y = df['target']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    random_state=42)

# Create a Decision Tree classifier
dt_classifier = DecisionTreeClassifier(random_state=42)

# Perform GridSearchCV to find the best hyperparameters
param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

grid_search = GridSearchCV(dt_classifier, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Get the best hyperparameters
best_dt_model = grid_search.best_estimator_

# Fit the best model to the training data
best_dt_model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = best_dt_model.predict(X_test)

# Calculate accuracy
test_accuracy = accuracy_score(y_test, y_pred)

# Print the results
print("Best Hyperparameters:", grid_search.best_params_)
print("Test Accuracy:", test_accuracy)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

# Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Best Hyperparameters: {'criterion': 'gini', 'max\_depth': None,

```
'min_samples_leaf': 1, 'min_samples_split': 2}
Test Accuracy: 0.9766536964980544
Confusion Matrix:
[[132   0]
 [  6 119]]
Classification Report:
              precision    recall  f1-score   support

     0       0.96       1.00       0.98       132
     1       1.00       0.95       0.98       125

 accuracy          0.98
 macro avg         0.98
weighted avg         0.98
```

## 11 Decision Tree classifier with Gini impurity and Entropy

```
[42]: import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report

# Load the dataset and perform train-test split (replace with your data loading
    process)
# Assuming 'df' contains the dataset with the features and target 'target'
# For example:
# df = pd.read_csv('your_dataset.csv')

X = df.drop(columns=['target'])
y = df['target']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, \
    random_state=42)

# Create a Decision Tree classifier with Gini impurity
dt_gini_classifier = DecisionTreeClassifier(criterion='gini', random_state=42)

# Perform GridSearchCV to find the best hyperparameters for Gini impurity
param_grid = {
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
```



```

grid_search_gini = GridSearchCV(dt_gini_classifier, param_grid, cv=5)
grid_search_gini.fit(X_train, y_train)

# Get the best hyperparameters for Gini impurity
best_dt_gini_model = grid_search_gini.best_estimator_

# Fit the best model to the training data for Gini impurity
best_dt_gini_model.fit(X_train, y_train)

# Make predictions on the test data for Gini impurity
y_pred_gini = best_dt_gini_model.predict(X_test)

# Calculate accuracy for Gini impurity
test_accuracy_gini = accuracy_score(y_test, y_pred_gini)

# Confusion Matrix for Gini impurity
conf_matrix_gini = confusion_matrix(y_test, y_pred_gini)

# Create a Decision Tree classifier with entropy
dt_entropy_classifier = DecisionTreeClassifier(criterion='entropy',
    ↪random_state=42)

# Perform GridSearchCV to find the best hyperparameters for entropy
grid_search_entropy = GridSearchCV(dt_entropy_classifier, param_grid, cv=5)
grid_search_entropy.fit(X_train, y_train)

# Get the best hyperparameters for entropy
best_dt_entropy_model = grid_search_entropy.best_estimator_

# Fit the best model to the training data for entropy
best_dt_entropy_model.fit(X_train, y_train)

# Make predictions on the test data for entropy
y_pred_entropy = best_dt_entropy_model.predict(X_test)

# Calculate accuracy for entropy
test_accuracy_entropy = accuracy_score(y_test, y_pred_entropy)

# Confusion Matrix for entropy
conf_matrix_entropy = confusion_matrix(y_test, y_pred_entropy)

# Print the results for Gini impurity
print("Gini Impurity - Best Hyperparameters:", grid_search_gini.best_params_)
print("Gini Impurity - Test Accuracy:", test_accuracy_gini)
print("Gini Impurity - Confusion Matrix:")
print(conf_matrix_gini)

```

```
# Print the results for entropy
print("Entropy - Best Hyperparameters:", grid_search_entropy.best_params_)
print("Entropy - Test Accuracy:", test_accuracy_entropy)
print("Entropy - Confusion Matrix:")
print(conf_matrix_entropy)
```

```
Gini Impurity - Best Hyperparameters: {'max_depth': None, 'min_samples_leaf': 1,
'min_samples_split': 2}
Gini Impurity - Test Accuracy: 0.9766536964980544
Gini Impurity - Confusion Matrix:
[[132   0]
 [  6 119]]
Entropy - Best Hyperparameters: {'max_depth': None, 'min_samples_leaf': 1,
'min_samples_split': 2}
Entropy - Test Accuracy: 0.9883268482490273
Entropy - Confusion Matrix:
[[132   0]
 [  3 122]]
```

```
[43]: pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in
c:\users\gayathriboddu\anaconda3\lib\site-packages (0.11.0)
Requirement already satisfied: scipy>=1.5.0 in
c:\users\gayathriboddu\anaconda3\lib\site-packages (from imbalanced-learn)
(1.7.3)
Requirement already satisfied: numpy>=1.17.3 in
c:\users\gayathriboddu\anaconda3\lib\site-packages (from imbalanced-learn)
(1.21.5)
Requirement already satisfied: scikit-learn>=1.0.2 in
c:\users\gayathriboddu\anaconda3\lib\site-packages (from imbalanced-learn)
(1.0.2)
Requirement already satisfied: joblib>=1.1.1 in
c:\users\gayathriboddu\anaconda3\lib\site-packages (from imbalanced-learn)
(1.3.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\users\gayathriboddu\anaconda3\lib\site-packages (from imbalanced-learn)
(2.2.0)
Note: you may need to restart the kernel to use updated packages.
```

## 12 Data Augmentation

```
[44]: from sklearn.model_selection import cross_val_score
from imblearn.over_sampling import SMOTE

# Data Augmentation with SMOTE
```

```

smote = SMOTE(random_state=42)
X_train_augmented, y_train_augmented = smote.fit_resample(X_train, y_train)

# GridSearchCV with augmented data
grid_search.fit(X_train_augmented, y_train_augmented)

# Get the best hyperparameters
best_model = grid_search.best_estimator_
best_model.fit(X_train_augmented, y_train_augmented)

# Calculate train accuracy and test accuracy
train_accuracy = best_model.score(X_train, y_train)
test_accuracy = best_model.score(X_test, y_test)

# Cross-validation for validation accuracy
val_accuracy = cross_val_score(best_model, X_train_augmented,
    ↪y_train_augmented, cv=5).mean()

print("Train Accuracy:", train_accuracy)
print("Test Accuracy:", test_accuracy)
print("Validation Accuracy:", val_accuracy)

y_pred = best_model.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

```

```

Train Accuracy: 1.0
Test Accuracy: 0.9766536964980544
Validation Accuracy: 0.9750543478260869
Confusion Matrix:
[[132   0]
 [  6 119]]

```

```

[45]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,
    ↪classification_report

# Load and explore the dataset:
# Replace 'path_to_csv' with the actual path to your downloaded CSV file.
data = df
# Print some information about the dataset.

```

```

print(data.head()) # Show the first few rows
print(data.info()) # Summary of the dataset

# Data Preprocessing:

# Drop rows with missing values
data.dropna(inplace=True)

# Split the data into features (X) and target (y)
X = data.drop('target', axis=1) # Assuming 'target' is the target column
y = data['target']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Scale the features to have zero mean and unit variance
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train a Gradient Boosting Classifier model:
# Create the Gradient Boosting Classifier model
model = GradientBoostingClassifier(n_estimators=100, random_state=42)

# Train the model
model.fit(X_train_scaled, y_train)

# Make Predictions:
# Predict on the test set
y_pred = model.predict(X_test_scaled)

# Evaluate the Model:

# Calculate accuracy and other metrics
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)

print("Accuracy:", accuracy)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", classification_rep)

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	

3	61	1	0	148	203	0	1	161	0	0.0	2
4	62	0	0	138	294	1	1	106	0	1.9	1

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1025 entries, 0 to 1024

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	age	1025 non-null	int64
1	sex	1025 non-null	int64
2	cp	1025 non-null	int64
3	trestbps	1025 non-null	int64
4	chol	1025 non-null	int64
5	fbs	1025 non-null	int64
6	restecg	1025 non-null	int64
7	thalach	1025 non-null	int64
8	exang	1025 non-null	int64
9	oldpeak	1025 non-null	float64
10	slope	1025 non-null	int64
11	ca	1025 non-null	int64
12	thal	1025 non-null	int64
13	target	1025 non-null	int64

dtypes: float64(1), int64(13)

memory usage: 112.2 KB

None

Accuracy: 0.9317073170731708

Confusion Matrix:

```
[[93  9]
 [ 5 98]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.91	0.93	102
1	0.92	0.95	0.93	103
accuracy			0.93	205
macro avg	0.93	0.93	0.93	205
weighted avg	0.93	0.93	0.93	205

[ ]:

[ ]: